

PLEASE JOIN HERE FOR FREE MATERIALS

LINK FOR TELEGRAM: [CLICK HERE](#)

LINK FOR INSTAGRAM : [CLICK HERE](#)

LINK FOR PAID MATERIALS: [CLICK HERE](#)

02-02-2-23

Kubernetes (k8s)

Kubernetes is an container cluster manager used for managing, monitoring, scaling the containerized applications on a clustered environment.

Kubernetes is used for deploying and running microservices based applications, to understand the crux and the way kubernetes works, it would good to know the basics of microservices application and their architecture

There are 2 architectures in which an application can be build

1. Monolithic application architecture
2. Microservices application architecture

#1. Monolithic Application architecture

The Enterprise application which consists of several modules are built into one single sourcecode project and will be build and packaged into one single deployable artifact is called "Monolithic application architecture"

There are few characteristics of Monolithic application:

1. the entire application is built into one single sourcecode, that includes all the modules or functional areas
2. A team of developers will be working across the functional areas or modules of the project
3. The whole application will be managed in single sourcecode management repository
4. The build system will produce an single deployable artifact
5. the whole application will be deployed and ran on a single enterprise application server environment

advantages:-

1. everyone in the team knows all the modules/functional areas of the system, since everyone works on end-end of the project
2. easy to build, package and deploy the application
3. scalability can be achieved easily

dis-advantages:-

1. since is an enterprise large application that comprises of lot of functionalities or modules and it is build into single sourcecode, the developers often find it very complex to work with entire system.
2. since the whole application is built into one single project
 - 2.1 To build the project it takes lot of time
 - 2.2 IDE are overloaded with lot of sourcecode files
 - 2.3 application servers takes more time in starting up

=====

03-2-2-2023

Kubernetes

Kubernetes is an container cluster manager, that takes care of running, monitoring and managing the containers on the cluster network of machines

There are 2 architectures on which applications are build

1. Monolithic application architecture
2. Microservices application architecture

1. Monolithic Application Architecture

An Enterprise large application that comprises of multiple

functional areas/modules are build together into one single sourcecode and is packaged into an single deployable artifact is called "Monolithic application"

Characteristics:

1. The enterprise application with all the modules are functionalities are built into one single sourcecode project
2. The Team of developers will working across all the modules of the project
3. A single sourcecode repository will be used for versioning and collaborating the development
4. only one single deployable artifact will be produced out of the build process
5. The application will be deployed on enterprise application server

advantages:

1. anyone in the project can work on any of the modules of the system
2. easy to build, package and deploy the application
3. scalability is easy to achieve in monolithic applications

dis-advantages:-

1. since the application is very huge and is build into one single sourcecode, developers often find it very complex to understand and work on the entire system
2. because of huge sourcecode, the integrated development environments (ide) will be overloaded and quickly goes un-responsive
3. building and packaging of the application becomes heavy and takes more time, due to which development will be impacted
4. application servers takes more time in deploying and starting the application, because of large in nature, due to which the productivity of the developer will be degraded

#2. Microservices application architecture

The Enterprise large application that comprises of several functional areas or modules are broken down into smaller services

=====

04-02-2023

Microservices Application Architecture

The Enterprise large scale application that comprises of multiple functional areas or modules is broken down into smaller microservice applications which are independently deployable and loosely coupled applications are called "Microservices"

Characteristics:-

1. Each Microservice application is built into its own sourcecode project which are independent of other modules/services of the system
2. Each Microservice application/project has its own sourcecode repository
3. Each Microservice has its own database schema into which those performs persistence operations
4. Each Microservice application is built into an independently deployable artifact, that is deployed on its own server runtime
5. Each Microservice is build by a team independent from other teams

=====

05-02-2023

Kubernetes (k8s)

Kubernetes is an container cluster manager, it takes care of distributing containerized applications on a cluster of computers, scheduling, monitoring and scaling up the containerized applications

while we are working on containerized applications, running them on a cluster has lot of problems:

1. distributing the containerized applications on a network of computers is very difficult
2. Keeping tracking of which containers are running on which nodes of the cluster is very difficult
3. patching and upgrading the containers are very difficult, we need to identify on which nodes of the cluster these containerized applications are running, we need to manually ssh onto those nodes and stop the running containers and rollout the newer versions of them.
and we need to determine the upgrade strategy to avoid downtime
4. monitoring and replacing an container in case of crash is very difficult
5. keeping track of the resource capacity and their utilization, and incase if a container is consuming huge system capacity on a node, which doesnt have sufficient resources, then migrating the container to another node that has enough capacity is very difficult
6. scaling up the applications on network cluster is difficult

To overcome the above problems/difficulties in managing and running the containerized applications on a network we need kubernetes.

Kubernetes Architecture

Kubernetes is an container cluster manager that takes care of scheduling, monitoring and managing the containerized applications on a network cluster of machines

The Kubernetes has 4 major components in it

1. Master Node
2. Worker Node
3. Kubectl
4. etcd

=====

06-02-2023

Kubernetes (k8s)

Kubernetes is an container cluster manager that takes care of scheduling, monitoring and managing the containers over the cluster of machines.

There are 4 major components are there in kubernetes

1. Master Node or Control Plane
2. Worker Node
3. Kubectl
4. etcd

#1. Kubernetes Master Node

Kubernetes Master is the central component of the kubernetes cluster, he is even called as "Control Plane", the job of scheduling the pods, monitoring and managing them across the worker nodes of the cluster will be taken care by "Kubernetes Master".

There are 3 sub-components are there in MasterNode

- 1.1 Api Manager
- 1.2 Scheduler
- 1.3 Controller Manager

1.1 Api Manager

The Api manager is the http endpoint which acts as an face or front-end to the kubernetes cluster/master. To perform any action or an operation on the kubernetes cluster, the clients or we need to talk through the api manager only.

The Api manager is an endpoint built on http protocol and exposes itself to the world letting them interact with Kubernetes Master.

There are many ways we can talk to api manager

1. The Kubernetes has provided an CLI tool called "kubectl" using which we can talk to the api manager
2. we can invoke the httpendpoints exposed by the api manager directly
3. there are api libraries provided by the kubernetes itself, using them we can connect and interact with api manager

The api manager performs 2 major activities

1. the api manager upon receiving the request will authenticate whether user who is sending the request is an valid user or not
and checks whether the user is authorized to perform the requested operation or not

2. in addition to the above, it validates the kubernetes spec that is send aspart of the request before passing it to the scheduler
for performing further operation

1.2 Scheduler

upon receiving and validating the request, the api manager to further process the request it handovers the request to the Scheduler.

The Scheduler is responsible for scheduling an pod for execution on the worker nodes of the cluster. The Scheduler will talks to the kubelet process that is running on each worker node of the cluster, checks does the workernode has enough capacity available for running the pod or not, if not available goes to the next worker node of the cluster until it finds one

upon identifying an workernode with enough capacity, the scheduler will handovers the request to the kubelet process of the workernode asking to bringup the pod on that node by allocating the requested resources

1.3 Controller Manager

The Controller Manager is an daemon/background process that is running on the kubernetes manager/control plane and ensures to bring the cluster to the desired state. It continous monitors and interacts with the kubelet process of the workernodes to identify the system is in desired state or not

There are 5 types of controllers are there

1. ReplicaSet

ReplicaSet controller ensures the desired number of replicas of a pod is distributed across different nodes of the cluster.

In the event of a crash, the ReplicaSet controller replaces that with an running pod on the cluster.

The ReplicaSet controller ensures the desired number of replicas are always running the cluster

2. DaemonSet

DamonSet ensures a pod is running on all the nodes of the cluster. Incase if a new workernode has been added to cluster,

the daemonset ensures the pod is broughtup onto the new worker node aswell.

3. DeploymentSet

DeploymentSet controller helps us in upgrading or patching the old pods with newer versions by supporting different deployment strategies

4. Service

service is an controller that discovers the running pods on the nodes of the cluster based on labels and register

with them, so that it can distribute the request to the pod application by

loadbalancing.

5. Job

Job controller helps us in running a job or script on a node of the cluster to perform one-time operation

=====

=====

07-02-2023

Kubernetes (k8s)

Kubernetes is an container cluster manager that takes care of scheduling, monitoring and managing the containers on the cluster of computers.

There are 4 components are there in kubernetes

1. master or control plane
2. workernode
3. kubectl
4. etcd

1. Master Node or Control Plane

it is the central component of the kubernetes cluster that takes care of scheduling, managing and monitoring the pods across the workernodes of the cluster

There are 3 more components are there in Master Node

1.1 Api Manager

acts as a front-end or interface in communicating with the master or control plane, api manager is built as an

rest/http endpoint through which people can interact with kubernetes master.

There are 2 major activities api manager performs

- 1.1.1 authentication and authorization
- 1.1.2 validating the request (specfile)

1.2 Scheduler

upon receiving the request, the api manager forwards the request to scheduler. The scheduler does the job of

communicating with the kubelet process of the workernode across the cluster to identify an workernode suitable

enough in running the pod and hands over the job of bringing up the pod on the workernode to the kubelet process

1.3 Controller Manager

The Controller Manager is a daemon process, that always ensures for bringing the cluster to the desired state.

There are 5 types of controllers are there

1.3.1 ReplicaSet Controller = desired no of replicas of a pod are running on the cluster will be taken care by ReplicaSet controller

1.3.2 DaemonSet Controller = ensures a pod is running always across all the nodes of the cluster

1.3.3 Service Controller = discovers the running pods on the cluster based on the

labels and groups them and expose them over the cluster, it takes care of loadbalancing the request received across the pods registered with that service

1.3.4 DeploymentSet Controller = used for upgrading, patching the pods based on upgrade/rollout strategies

1.3.5 Job Controller = helps us in running one-time operation on the nodes of the cluster

#2. Worker Node

Kubernetes Worker Node is an physical server/virtual machine/cloud instance that is attached to the kubernetes

master/control plane on which the pods are scheduled for execution. On a workernode there could be one or more pods be scheduled for execution based on the capacity of the worker node

There are 3 major components to be installed on the each workernode in order to register and schedule pods on them

1. container runtime

A workernode must be installed with an containerization engine, so that the containerized applications can be executed on them.

For eg.. we can install docker on the workernodes to run docker containers on them

2. kubelet

kubelet is an process that runs on each workernode of the cluster. The Kubelet process acts as an agent,

in letting the controlplane or master interact with the Workernodes of the cluster. The Master Node schedules a pod for execution on the cluster by handovering the job to kubelet process only.

Additionally Kubelet process performs various activities like

1. kubelet gathers the information about the running pods and their status and reports them to the control plane or master whenever requested

2. The job of running the pods and bringing them up on a workernode is taken care by kubelet process only

3. through the help of kubelet process, the control plane can determine which workernode is suitable for running or scheduling a pod for execution

3. kubeproxy

kubeproxy enables the traffic to the external network for a pod in the cluster

#3. Kubectl

Kubectl is an cli tool provided by the kubernetes through which we can communicate with the controlplane or master.

it helps us in administering, monitoring and managing the kubernetes cluster

#4 etcd

etcd is an key/value pair database where all the kubernetes objects information will be stored on the etcd database only

=====

08-02-2023

There are multiple ways we can setup an Kubernetes Cluster

1. on-premise with physical server machine
2. through virtual machines
3. Minikube
4. AWS EKS cluster

#4. AWS EKS Cluster

To work with AWS EKS Cluster we need to have an AWS Cloud Account, if not create one.

The aws cloudplatform has provided an service called EKS stands for "Elastic Kubernetes service",

it is an managed service provided by aws platform to host kubernetes on aws.

The AWS Cloudplatform itself takes care of provisioning the master/control plane, workernodes and

setup the cluster with CNI network and install necessary components like containerization engine, kubelet, kubeproxy etc

There are lot of advantages of using EKS Cluster over the one-premise

1. we dont need to manually setup, configure the kubernetes cluster, rather with an click of a button

the AWS itself takes care of provisioning the kubernetes for us

2. Monitoring the kubernetes cluster and keeping track of the health and incase if any of workernodes

are crashed those are replaced with the healthier nodes will be taken care by the aws itself

3. The high availability of the cluster will be guaranteed since the workernodes are distributed across

the availability zones of the vpc, in addition we can take the advantages of making the application accessed to customer with low-network latency

4. The AWS Cloud platform itself takes care scale-out/scale-in the cluster capacity/size based on the

load of the cluster so that we never run out of cluster capacity

5. The HA of the master/control plane will be taken care by the aws cloud platform itself

How many ways we can provision an EKS Cluster on AWS Cloud?

There are 3 ways we can setup EKS Cluster on AWS Cloud

1. aws management console = we can manually configure and setup the EKS Cluster through console
2. ekscli interface = ekscli is an command-line tool provided by aws team, using which we can quickly setup the eks cluster on the cloud
3. through automation using terraform or ansible

Let us understand how to setup an EKS Cluster

#1. we need an vpc "hondaekscluster"

vpcname: "hondaekscluster"

cidr: 10.0.0.0/16

#2. now we need to create the subnets

There are 3 ways we can host the eks cluster

1. public = we can create the worker nodes on public subnet of the vpc = only for development purposes we can use but should avoided

2. private & public = eks masternode / controlplane will be provisioned on public subnet and the workernodes are provisioned on private subnet.

Typically used in organization env, where teams should able to access the kubernetes cluster for deploying and running the applications on the cluster (non-production env)

3. private subnet = both master/workernodes should be provisioned on private subnet only (recommended for production usage).

To manage the cluster we need to use jumpbox

#3. to provision an eks cluster we need atleast 2 subnets under 2 different availability zones of an vpc (HA)

#4. when we are provisioning the workernodes under private subnet of an vpc, workernodes will not have access to the public network,

so they cannot pull the docker images from dockerhub to run pods on the container, so we need to must and should setup ecr (elastic container registry)

provided by aws cloudplatform, and publish the docker images on the ecr registry, so that the workernodes can pull the docker images from aws ecr.

=====

09-02-2023

#1. setup the workstation

1. install openjdk-11-jdk

sudo apt update -y

sudo apt install -y openjdk-11-jdk

2. install apache maven

sudo apt install -y maven

3. install vscode

download the vscode binary ".deb" from the vscode downloads, by default it will be downloaded into ~/Downloads directory

cd ~/Downloads

sudo apt install -f ./code_....deb

4. docker

4.1

sudo apt install -y ca-certificates curl gnupg lsb-release

4.2 download and add the gpg docker key

sudo mkdir -m 0755 -p /etc/apt/keyrings

curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /etc/apt/keyrings/docker.gpg

4.3 add the docker registry into ubuntu sources.list.d
echo \
"deb [arch=\$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg]
https://download.docker.com/linux/ubuntu \
\$(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list >
/dev/null

4.4 update the registry
sudo apt update -y

4.5 install docker
sudo apt install -y docker-ce

4.6 grant sriman access to the docker by adding him to the docker group
sudo usermod -aG docker \$USER
exit the terminal and re-enter or restart

run the below command to verify sriman has access to docker
docker image ls

5. create the IAM User on aws cloud account

5.1. group: hondaopsgroup
attached policies:
1. AmazonEC2FullAccess
2. AmazonEKSClusterPolicy
3. AmazonEKSWorkerNodePolicy
4. AmazonEC2ContainerRegistryFullAccess
5. AmazonEKSServicePolicy
6. AmazonEKS_CNI_Policy

5.2 create a user and add him to the above group

username: bob
add: hondaopsgroup

5.3 generate the api keys for accessing
goto user, click on security credentials and click on generate accesskeys, once
generated download them

5. install AWScli on workstation

<https://docs.aws.amazon.com/cli/latest/userguide/getting-started-install.html>
curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o "awscliv2.zip"
unzip awscliv2.zip
sudo ./aws/install

6. create aws credentials using awscli tool to access the cloud account resources
using cli

aws configure
prompts for
api key:
access key:
region: ap-south-1
output format: none

=====

10-02-2023

#2. How to setup the Elastic Container Registry?

2.1

There are 2 types of Elastic Container Repositories are there

1. private repository = only the IAM user who has necessary policies attached only can access the repository
2. public repository = any one can access, pull/publish images into the repository

2.2

per each docker image we want to publish we need to create on repository in ecr

2.3 to publish or pull images from these repositories we need to login into ecr registry. and in case of private repository

we need have AmazonEC2ContainerRegistryFullAccess to login or pull or push images

2.4 goto IAM User we have setup earlier and attach a policy to either user level or group level with policy: AmazonEC2ContainerRegistryFullAccess

2.5 goto elastic container registry and click on Get Starter or create repository

1. private repository and enter an name for repository
2. go into the repository and click on View Push Commands for login pull or push instructions

```
aws ecr get-login-password --region <regionName> | docker login --username AWS  
--password-stdin registryURL
```

How to provision the EKS Cluster?

1.

=====

12-02-2023

How to provision an eks cluster?

There are 2 parts are there in setting up an eks cluster on aws cloud platform

1. control plane or master node
2. worker nodes

There are 3 topologies where we can choose one of them in setting up the eks cluster

1. both master/workernodes on public subnets = not recommended since the whole cluster is exposed to the world and poses security

2. master on public subnet and workernodes on private subnet = usually used in organizations
allowing the team of people to manage the cluster directly (not recommended for production usage)
3. both master/workernodes on private subnet only = highly recommended for production usage

we wanted to setup the eks cluster using option-#2

#1. create the vpc
vpcname: hondaeksvpc
cidr: 10.0.0.0/16

#2. create 2 public subnets and 2 private subnets, public subnets for control plane and private subnets for workernodes
we need to create 2 of them for each public/private subnets across the azs of the region, to ensure high-availability
2.1 hondaekspubsn1, 10.0.1.0/24
2.2 hondaekspubsn2, 10.0.2.0/24
2.3 hondaeksprvsn3, 10.0.3.0/24
2.4 hondaeksprvsn4, 10.0.4.0/24

#3. provision internet gateway and attach to the vpc
internet gateway name: hondaeksig
attach to vpc: hondaeksvpc

#4. route the public network traffic through internet gateway by creating routetable
routetable name: hondaigrt
subnet association: hondaekspubsn1, hondaekspubsn2
route: 0.0.0.0/0 -> hondaeksig

provision eks control plane or master node

1. The eks master node or control plane takes care of provisioning, setting up and installing the workernodes based on the configuration,
to let the eks cluster manage the workernode we need to attach an IAM Role to the master node while provisioning. So let us first setup the IAM Role
1.1 goto IAM Policies service and choose create role
1.2 within the create role, select the category type: EKSCluster
1.3 add the policy as EKSClusterPolicy into the Role and create it

2. goto Elastic Kubernetes Service and click on create cluster.
2.1 cluster name
2.2 choose the vpc
2.3 choose the way you want to host the cluster (master/controlplane : public subnet, workernodes: private subnet)
2.4 Attach Role we created above
and click on create cluster

upon creating the cluster, it setup or provision only control plane without workernodes, now navigate into the cluster and click on create nodegroup

setting up workernodes

1. to setup the workernodes we need create an IAM Role that should be attached to the workernode during provision.

goto IAM Policies and choose Role and create new

choose the type: EKSNodeGroupRole

policies:

1. AmazonEKSWorkerNodePolicy
 2. AmazonEKSContainerRegistryReadOnly
 3. AmazonEKS_CNI_Policy
- create role

2. goto the eks cluster we have created above and click on add NodeGroup
we can think of a NodeGroup as equivalent to ASG. The NodeGroup takes care of provisioning the workernodes and attach to the EKS MasterNode.

here we can specify

1. shape of the workernode (t2.micro)
2. min, max, initial workernodes
3. scale-out threshold
4. subnets

The NodeGroup based on the above configuration takes care of provisioning and managing the workers automatically

How to setup the kubectl on ubuntu?

documentation:

<https://kubernetes.io/docs/tasks/tools/install-kubectl-linux/#install-using-native-package-management>

#1. install pre-requisite

1. sudo apt-get update
2. sudo apt-get install -y ca-certificates curl
3. sudo apt-get install -y apt-transport-https

#2. add gpgkey

```
sudo curl -fsSLo /etc/apt/keyrings/kubernetes-archive-keyring.gpg  
https://packages.cloud.google.com/apt/doc/apt-key.gpg
```

#3. now add the kubernetes repository to the ubuntu sources.d or sources.list

```
echo "deb [signed-by=/etc/apt/keyrings/kubernetes-archive-keyring.gpg]  
https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee  
/etc/apt/sources.list.d/kubernetes.list
```

#4. update the repository and install kubectl

```
sudo apt update  
sudo apt install -y kubectl
```

How to add the kubeconfig to workstation to access the eks cluster using kubectl run the below command.

```
aws eks update-kubeconfig --region <REGION> --name <ClusterName>
aws eks update-kubeconfig --region ap-south-1 --name HondaCluster
```

How to install kubectl on windows 10/11?

we just need to download the kubectl.exe and place it in the c:\minikube directory
1. goto c:\minikube\ directory, if we dont have create one manually and run the below command on powershell/command-prompt

```
curl.exe -LO "https://dl.k8s.io/release/v1.26.0/bin/windows/amd64/kubectl.exe"
```

2. set the PATH variables pointing to the c:\minikube directory

3. kubectl version

```
C:\Users\Sriman>kubectl version
```

```
WARNING: This version information is deprecated and will be replaced with the output from kubectl version --short.
```

Use --output=yaml|json to get the full version.

```
Client Version: version.Info{Major:"1", Minor:"26", GitVersion:"v1.26.0",
GitCommit:"b46a3f887ca979b1a5d14fd39cb1af43e7e5d12d",
GitTreeState:"clean", BuildDate:"2022-12-08T19:58:30Z", GoVersion:"go1.19.4",
Compiler:"gc", Platform:"windows/amd64"}
```

```
Kustomize Version: v4.5.7
```

```
Server Version: version.Info{Major:"1", Minor:"26", GitVersion:"v1.26.1",
GitCommit:"8f94681cd294aa8cfd3407b8191f6c70214973a4",
GitTreeState:"clean", BuildDate:"2023-01-18T15:51:25Z", GoVersion:"go1.19.5",
Compiler:"gc", Platform:"linux/amd64"}
```

How to install minikube on windows 10/11 machine?

Running an kubernetes cluster on a physical server environment or in a virtual machine env requires huge system resources and would take

lot of time in setting up on cluster. For local development and experimentation/learning we can use minikube installed

Minikube is an down-sized version of kubernetes cluster provided by kubernetes team, that takes care of installing and

running the kubernetes cluster on virtual machine with onesingle node/virtualmachine as master/workernode itself.

So that we can quickly run the cluster locally and work on local development

You can refer to the minikube start documentation:

<https://minikube.sigs.k8s.io/docs/start/>

before beginning installation of minikube, download and setup kubectl on the windows 10/11 machine

1. run the below command in powershell or windows command-prompt that downloads the minikube.exe and places under c:\minikube directory

```
New-Item -Path 'c:\' -Name 'minikube' -ItemType Directory -Force
Invoke-WebRequest -OutFile 'c:\minikube\minikube.exe' -Uri
```

'https://github.com/kubernetes/minikube/releases/latest/download/minikube-windows-amd64.exe' -UseBasicParsing

2. then add c:\minikube\ directory to the system path in environment variables.
exit the terminal and reopen

3. minikube start --driver=virtualbox

to verify the minikube cluster status run the below command
kubectl get po -A

```
C:\Users\Sriman>kubectl get po -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS
AGE				
kube-system	coredns-787d4945fb-zm2ch	1/1	Running	1 (25m ago)
28m				
kube-system	etcd-minikube	1/1	Running	1 (25m ago)
28m				
kube-system	kube-apiserver-minikube	1/1	Running	1 (25m ago)
28m				
kube-system	kube-controller-manager-minikube	1/1	Running	1 (25m ago)
29m				
kube-system	kube-proxy-znldh	1/1	Running	1 (25m ago)
28m				
kube-system	kube-scheduler-minikube	1/1	Running	1 (25m ago)
28m				
kube-system	storage-provisioner	1/1	Running	1 (25m ago)
28m				

Kubernetes Dashboard

Kubernetes has provided an management-console or a kubernetes cluster dashboard using which we can access all the kubernetes object through the web console like

- namespaces
 - pods
 - services
 - deployments
- etc

we can run the kubernetes dashboard on minikube using
minikube dashboard

this installs/configures the dashboard on minikube and opens the console on the browser.....

.....
.....
13-02-2023

Kubernetes Namespace

Namespaces are used for creating naming compartments or logical grouping of objects/resources on the kubernetes cluster. A kubernetes cluster would be shared across multiple teams to run various different projects on the cluster, so to avoid objects/resources of one project/team are accessed by others we use kubernetes Namespaces

A kubernetes administrator can create multiple kubernetes users, associate them into groups. create an kubernetes namespace and grant access to the users/groups to those namespaces only.

In such a way per project/application/team, the administrator can create namespace associate users/groups into that namespace allowing them to run, manage modify the objects only within that namespace, so that we can avoid different people of different groups to see or manage resources of others, hence the cluster would be secured and protected.

By default aspart of the kubernetes install there are 4 namespaces are created

1. default = every object that is created in kubernetes cluster will be placed by default under "default" namespace only. The default namespace is by default empty. all the users/groups of kubernetes cluster has access to the default namespace and should be sufficient for most of the usecases
2. kube-system = all the kubernetes system objects like api manager, scheduler and controller manager etc are placed under "kube-system" namespace only
3. kube-public = by default kube-public ns is empty, if we place any objects within the kube-public, those are accessible publicly to everyone without authentication
4. kube-node-lease = all the kubernetes leased objects are associated to this namespace

1. How to see the namespaces that are on the cluster?

kubectl get (namespaces or ns)

returns all the list of namespaces that are available on the cluster

2. How to create our own namespace in the kubernetes cluster?

kubectl create namespace namespaceName

3. how to delete a namespace from the cluster?

kubectl delete namespace namespaceName

4. How to see all the pods running/available on the cluster?

kubectl get pods

by default the kubectl will be running on "default" namespace so it shows all the objects like pods, deploymentset, replicaset etc only in default namespace

if we want to see the objects/resources under a specific namespace then we need to use

kubectl get pods -n namespace

5. How to see all the objects on the cluster irrespective of the namespace?

`kubect1 get pods -A`

14-02-2023

What is a namespace, what is the purpose of it in kubernetes cluster?

Namespaces are used for creating logical grouping of objects & resources within the kubernetes cluster, so that we can isolate the objects & resources belongs to one project or domain or business unit from other by sharing the same cluster across multiple teams/business units.

Kubernetes administrator can create users, groups associate them to an namespace granting access, so that users can access/manage resources/objects of that namespace only securing the cluster

There are 4 namespaces will be created by default aspart of kubernetes install

1. default = by default, default namespace is empty. all the objects/resources that are created would be placed by default under default namespace
2. kube-system = The kubernetes system objects like api manager, controller manager and scheduler are created under kube-system namespace
3. kube-public = by default empty, and the objects/resources placed here are publicly accessible without authentication
4. kube-node-lease = leased objects are placed in kube-node-lease

1. How to see all the namespaces within the cluster?

`kubect1 get ns or namespaces`

2. how to create a new namespace?

`kubect1 create namespace namespacename`

3. how to see objects/resources within a namespace?

`kubect1 get pods -n namespacename`

4. how to delete a namespace?

`kubect1 delete namespace namespacename`

5. how to see objects/resources from all the namespaces of the cluster?

`kubect1 get pods -A`

How does the kubect1 will be configured to access the kubernetes cluster?

during the time of kubernetes cluster install, the kubect1 will be configured to point to the cluster. Incase of AWS EKS cluster

the aws cli has provided an cli command through which we can setup kubect1 pointing to the AWS EKS cluster as

`aws eks update-config --region regionName --cluster clusterName`

The kubernetes cluster information will be stored by default under \$HOME/.kube/config file. It has majorly 3 sections in it

1. clusters = clusters section lists all the kubernetes cluster and their certificate key using which it needs to communicate with the kubernetes cluster incase of https

2. users = all the users of different clusters using which we need to connect to the cluster

3. contexts = a context is a name given for the combination of cluster, user and namespace which would be specified to kubectl to be used in connecting and managing the cluster

\$HOME/.kube/config

```
-----
apiVersion: v1
kind: config
preferences: {}
clusters:
  - cluster:
      name: testcluster
      server: http://host:port
      certificate-authority-data: key
  - cluster:
      name: stagecluster
      server: http://host:port
      certificate-authority-data: key
users:
  - name: testclusteruser
    user:
      client-certificate-data:
      client-key-data:
  - name: stageclusteruser
    user:
      client-certificate-data:
      client-key-data:
contexts:
  - context:
      name: testcontext
      cluster: testcluster
      user: testclusteruser
      namespace: default
  - context:
      name: stagecontext
      cluster: stagecluster
      user: stageclusteruser
      namespace: ns1
current-context: testcontext
```

How to switch from one cluster to another cluster?
we need to change the current-context attribute in kube config file

There are 2 ways to change it

1. we can directly goto .kube/config file and modify it using text editor, which is not recommended
2. we can modify using kubectl cli commands

1. how to switch to an another context using kubectl?
kubectl config use-context stagecontext

2. how to see all the contexts in kubeconfig?
kubectl config get-contexts

3. how to see the current-context in use?
kubectl config current-context

4. how to change a namespace of a kubernetes cluster
kubectl config set-context --current --namespace=ns1

15-02-2023

How can we connect or manage multiple kubernetes cluster using kubectl?
the cluster information would be by default configured or stored in kubernetes workstation at \$HOME/.kube/config. The kubeconfig file contains 3 parts

1. clusters = In the clusters section information about the kubernetes cluster and the SSL keys will be available
2. users = The list of users pertaining to each cluster using whom we want to connect to the cluster along with secretkeys will be available
3. contexts = is a combination of cluster + user with namespace that should be used by kubectl to connect

.kube/config

```
apiVersion: v1
kind: config
preferences: {}
clusters:
  - cluster:
      name: cluster1
      server: http://host:port
      certificate-authority-data:
  - cluster:
      name: cluster2
      server: http://host:port
      certificate-authority-data:
users:
```

```

- name: clusteruser1
  user:
    client-certificate-data:
    client-key-data:
- name: clusteruser2
  user:
    client-certificate-data:
    client-key-data:
contexts:
- context:
  name: devcontext
  cluster: cluster1
  user: clusteruser1
  namespace: ns1
- context:
  name: testcontext
  cluster: cluster2
  user: clusteruser2
  namespace: ns2
current-context: devcontext

```

How to switch from one cluster to another cluster?

To switch from one cluster to another we need to modify current-context attribute in .kube/config file. There are 2 ways we can modify the .kube/config file

1. manually we can edit the file using text editor
2. use kubectl cli to modify the .kube/config file

1. How to see the current-context in which we are in?

kubectl config current-context

2. how to see the kubeconfig file

kubectl config view

(or)

cat ~/.kube/config

3. how to get all the contexts configured?

kubectl config get-contexts

4. How to switch from one context to a different context?

kubectl config use-context contextName

5. how to change the namespace in a context?

kubectl config --set-context --current --namespace=namespaceName

How to add a new kubernetes cluster information in .kube/config file?

There are 2 ways we can add cluster information into .kube/config file

1. modify the .kube/config file manually and add the cluster information
2. use --set-cluster in kubectl

kubectl config --kubeconfig=config set-cluster development --server

http://host:port --insecure-skip-tls-verify

Kubernetes Objects

There are different types of resources are there in kubernetes

1. Pod
2. ReplicaSet
3. DeploymentSet
4. DaemonSet
5. Service
6. Job
7. Ingress
8. Loadbalancer

The kubernetes will represent all these resources information in terms of objects and persist them in etc of the kubernetes cluster. all these objects are created under kube-system namespace. The state of the kubernetes cluster will be represented by the type/number of objects on the cluster

A kubernetes object is a record of intent, once we created a kubernetes object, the system constantly works to ensure the object exists or is created on the cluster. Through the objects we are specifying the kubernetes system, what we want on the cluster

How to create the objects on the cluster?

These kubernetes objects can be created by writing resource spec or manifest file and should be passed as an input to the control plane using kubectl.

The Control Plane reads the resource spec or manifest and creates an appropriate kubernetes object and stores in etcd

1. write a resourcespec based on the type of the object/resource we want to create
2. by using kubectl pass it as an input to control plane /master node
3. master node validates the spec and reads/extracts the information and creates appropriate object and stores in etcd
4. the system constantly works to bring the desired state

The resource spec or manifest is a YAML file in which we describe the information about the object we want to create in terms of key/value pair and pass it as an input to control plane. The structure and contents of the spec file is defined by kubernetes itself for each type of resource.

But all of these spec files carry few common attributes irrespective of type of object is:

1. apiVersion = which version of the kubernetes object we are using for creating
2. kind = type of object
3. metadata = used for defining labels for the object
3. namespace = under which namespace the object should be created
4. spec = the desired state of the object

How to create and manage these objects in kubernetes cluster?

The kubectl has provided sophisticated ways of creating, managing these objects on kubernetes cluster. There are 3 ways we can create these objects on the cluster.

1. imperative commands
2. imperative object configuration
3. declarative object configuration

1. imperative commands:

kubectl has provided handful commands to which we can pass arguments in creating various different types of kubernetes objects on the cluster. this avoid writing an specfile manually in creating objects.

```
kubectl run apache2 --image=apache2:latest --port=80
```

advantages:-

1. we can quickly create an object on the cluster without writing any manifest file and test it

dis-advantage:-

1. we dont have any spec file in hand, based on the inputs/arguments we passed, the kubectl creates the yml onfly and passes it to the control plane. if we want to modify or reuse in recreating the object at later point of time we dont have the spec with us

16-02-2023

What are kubernetes objects?

kubernetes object is an record of intent, through which we can tell what we wanted to create on kubernetes cluster. upon creating an object kubernetes cluster will tries to bring the system to the desired state based on object definition. all the kubernetes objects are persisted permanently aspart of etcd.

There are different types of objects are there

1. pods
 2. replicaset
 3. deploymentset
 4. daemonset
 5. job
 6. service
 7. ingress
- etc

by writing an kubernetes resource specfile we tell the kubernetes to create an object on the cluster. The devops engineer or an kubernetes developer would write the resource specfile defining/describing the info about the object on workstation and using the help of kubectl and passes the resource spec/manifest to the control plane or kubernetes master. Kubernetes master validates the spec and creates an

kubernetes object representing the spec and stores in the etc

The resource spec are yaml files, and the structure of these yaml files are specific to the type of the object we want to create and standardized by the kubernetes. based on the structure (key/value pairs) defined by the kubernetes we need follow and write the spec file in order to get this validated by the cluster. There are few common attributes are there across all type of object specs which are defined as below.

1. apiVersion = defines the specfile version being used in defining the object
2. kind = type of object
3. metadata= used for attaching labels to the object for identification and retrieval
4. namespace = under which namespace object should be created
5. spec = specification of the object

How many ways are there in creating these objects on the cluster?

There are 3 ways are there

1. imperative commands
2. imperative object configuration
3. declarative object configuration

1. imperative commands

The kubectl has provided handful of commands taking arguments as input in creating various different kubernetes objects like pods, service, deployments etc. we don't need to write the resource specfile in creating these objects

```
kubectl run podName --image=imageName:tag --port=portNo
```

advantages:-

1. quickly create an object on the cluster

dis-advantage:-

2. since we don't have resource specfile we cannot reuse the object across the environments

2. imperative object configuration

In an imperative object configuration, for each type of resource we want to create, we need write an resource specfile describing the information about the object and passes it as an input to the kubectl asking him to perform the operation like create or delete.

apache2-pod.yaml = in this specfile we described the pod information

```
kubectl create -f apache2-pod.yaml
```

```
kubectl delete -f apache2-pod.yaml
```

3. declarative object configuration

In order to run an application on kubernetes cluster, they could be bunch of objects we need to create on the cluster. So defining each object we need to write

the respective resource specfile. upon writing these specfiles we need to create each object by running the spec keeping the order of dependencies. It takes lot of time and would be hard to memorize and execute the objects in the dependency order

Instead keep all the resources under one directory under the project like manifests

```
/ configs
airtelcare2
|-manifests
  |-pod.yml
  |-deployment.yml
  |-service.yml
```

pass the directory as an input to the kubectl asking to apply these manifests on the cluster

```
kubectl apply -f manifests/
kubectl delete -f manifests/
```

Let us try to understand how to create different kubernetes objects by writing the specfile

1. Pod

Pod is an smallest unit or entity in the kubernetes world in which one or more containers are kept together and executed. In general we may have multiple containerized applications that shares common dependencies like

1. network
2. resources file/mounts/volumes
3. lifecycle (start/stop)

In such case instead of having these 2 applications running as 2 independent containers, to reduce the effort of managing them togetherly easily kubernetes has introduced an logical container/unit called "pod", in which we can place one or more containers as an single unit and can manage them together like starting or stopping etc

How to write an resource specfile in running an pod on the cluster?

application binary -> Dockerfile -> image -> container registry -> pod specfile -> create/apply #cluster

apache2 image

Apache2 image is already published and available on the container registry, so we can quickly write pod specfile in running an container ontop of the image on the cluster as below.

assumption: the image is already available in the registry

```
apache2-pod.yml
apiVersion: v1
```

```

kind: Pod
metadata:
  name: apache2pod
spec:
  containers: #defining the info about the containers we want to run inside this
pod
    - name: apache2
      image: apache:latest
      ports:
        - containerPort: 80
          name: http
          protocol: tcp

```

kubectl create -f apache2-pod.yml

1. how to see the pods on the cluster

kubectl get pods -n namespace

2. how to create pods on the cluster

kubectl create -f pod.yml

3. how to see the information about the pod ?

kubectl describe pod podName

4. how to acces the logs of the running pod on the cluster?

kubectl logs podname

5. how to get more information about the pods on the cluster?

kubectl get pods -o wide -n namespace

6. how to expose a pod to the host using port-foward?

kubectl port-forward podName hostPort:containerPort

7. how to delete a pod?

kubectl delete -f pod.yml

or

kubectl delete podName

 17-08-2023

What is a Pod?

Pod is an smallest entity or an unit within the kubernetes cluster, in which we run one or more container together inside it. There can be few containers that are dependent on each other in terms of sharing

1. resources

2. file system / volumes

3. lifecycle

rather than having such containers managed independently we can place them inside one pod and manage it together easily.

How to write a pod manifest file in creating and running a pod?

apache2-pod.yml

apiVersion: v1

kind: Pod

metadata:

 name: apache2pod

namespace: default

spec:

 containers:

 - name: apache2

 image: ubuntu/apache2:latest

 ports:

 - name: apache2port

 containerPort: 80

 protocol: TCP

ghp_Y0Nu9d01JNAPqakggYExU2FCIePrgK1TE6mt

repositoryName/image:tag

techsriman/sailor:1.0

18-02-2023

Readiness Probe and Liveness Probe

Upon placing an request for executing a pod the kubernetes makes best effort of choosing an node on which the pod will scheduled for execution. during the time of bringing up the pod, kubernetes will read the podspec file and identifies the image on which the pod container should be created and checks to see is the container image is available locally.

If the container image is not available locally, then it pulls the image from container registry onto that workernode and creates the pod and runs it. the moment the pod has been started kubernetes reports the pod under running state

but the underlying application running inside the pod may not be started or while running the underlying application might go into unresponsive state due to resource availability, stuck threads etc. even then also the kubernetes reports the pod status as running and available. The kubernetes doesnt know the real/underlying status of the application that is running inside pod, due to which it would be able to identify such faulty applications

In general kubernetes takes care of replacing a pod that is not working or crashed, but in this scenario since kubernetes doesn't know the status of the underlying application, it would assume the pod is working and leave it unresponsive

To let the kubernetes identify such faulty or unresponsive applications running inside the pod and replace them we need to provide the health information about the application to the kubernetes through `readinessProbe` and `livenessProbe`

Kubernetes supports 2 types of checks to be performed on a pod application

1. `readinessProbe`

upon scheduling a pod for execution the kubernetes pulls the image and runs the pod and reports the status of the pod as running, but the underlying application that is running inside the pod container could be still under starting state and may not be ready for accessing or accepting the requests. even then also kubernetes assumes the application has been started and every request received for the application from the outside world would be routed to the application which will result in failure.

To avoid this problem kubernetes has introduced `readinessProbe`, using which kubernetes determines when the actual application is available for accessing. The developer of the application has to expose an http endpoint and configure the information about this endpoint in podspec configuration file, letting kubernetes understand, it has to access this endpoint to verify the readiness of the application.

upon bringing up the pod, the kubernetes will perform the readiness check by periodically hitting or accessing the readiness Application Endpoint we configured in spec file, until the application has reported its availability, kubernetes will not route the request to the pod instance of the application. upon the application reports its availability, kubernetes marks the pod ready for scheduling the request and would stop performing `readinessProbe`

2. `livenessProbe`

while the pod is running on the cluster, there could be a chance due to resource issues or application failures the application running inside the pod may become unresponsive. kubernetes is not aware of such unresponsive pods, so that it routes the incoming requests for the pod application to these unresponsive pods as well which will eventually lead to failures.

if the kubernetes can somehow identify such unresponsive running pods on the cluster, it can terminate them and recreate another pod on the cluster, which can be done through `livenessProbe`

the application developer has to expose an http endpoint, that can be used by the kubernetes to periodically verify the application is running/responsive or not. while writing the podspec file the kubernetes developer has to configure the `livenessProbe` information so that kubernetes can perform this check for us

The livenessProbe would be started only upon the readinessProbe has been reported as successful and continued to perform the checks until the pod has been reported as failed or manually terminated

We have an application urotaxi (java web application), in this application we exposed #2 endpoints (readiness/liveness). The url for accessing these endpoints are below

1. /actuator/health/readiness
2. /actuator/health/liveness

now while writing the podspec file in running the application as pod on kubernetes cluster, we need to configure both readinessProbe and livenessProbe configuration as below.

urotaxi-pod.yml

apiVersion: v1

kind: Pod

metadata:

 name: urotaxipod

 labels:

 version: 1.0

spec:

 containers:

 - name: urotaxi

 image: techsriman/urotaxi:1.0

 ports:

 - name: tomcatport

 containerPorts: 8089

 protocol: TCP

 readinessProbe:

 httpGet:

 path: /actuator/health/readiness

 port: 8080

 initialDelaySeconds: 5

 timeoutSeconds: 10

 failureThreshold: 3

 livenessProbe:

 httpGet:

 path: /actuator/health/liveness

 port: 8080

 initialDelaySeconds: 5

 timeoutSeconds: 10

 failureThreshold: 3

19-02-2023

Resource declarations

whenever an pod application is running on the kubernetes cluster, it is going to consume cpu/memory during execution. The amount of cpu/memory it consumes depends on various factors like

1. hardware capacity of the machine on which we are running the application
2. how much amount of user traffic is coming to the application
3. the amount of data the application is processing

How can we know the cpu/memory utilization of an application?

The performance testing team puts the application for evaluation for deriving the beanch mark metrics of the cpu/memory usage levels of the application. The team puts an virtual load on the application based on the average number of users always going to access the application and derives the cpu/memory consumption.

whenever we scheduled an pod for execution, the kubernetes scheduler takes care of identifying an appropriate worknode on the cluster which has sufficient cpu/memory resources available for running the pod and handovers the pod execution to the kubelet process of the node.

So to the kubernetes scheduler determine the right workernode to be used in running the pod, we need to define or declare the resource specification aspart of the pod spec file. we define the minimal requirements in running an pod based on which the workernode will be choosen, incase if the pod application is requesting more than the resources requested, the kubelet process will try to accomodate the resources if available on the worker node.

In case if the workernode doesnt have sufficient resource capacity, the pod will be terminated and would rescheduled to execute on a workernode which has appropriate capacity. so we need to define the resource specification in the pod spec for execution

```
roadster-pod.yml
apiVersion: v1
kind: Pod
metadata:
  name: roadsterpod
spec:
  containers:
    - name: roadster
      image: techsriman/roadster:1.0
      ports:
        - name: tomcatport
          containerPort: 8080
          protocol: TCP
      readinessProbe:
        httpGet:
          path: /roadster/actuator/health/readiness
```

```

        port: 8080
        initialDelaySeconds: 10
        timeoutSeconds: 10
        failureThreshold: 3
    livenessProbe:
        httpGet:
            path: /roadster/actuator/health/liveness
            port: 8080
        initialDelaySeconds: 10
        timeoutSeconds: 10
        failureThreshold: 3
    resources:
        requests:
            cpu: "500m"
            memory: "512Mi"
        limits:
            cpu: "1000m"
            memory: "1024Mi"

```


 20-02-2023

what are resource declarations, why do we need to use them?

Resource declarations are the way through which we specify the cpu and the memory requirements and limits for a pod to be scheduled and execute on the cluster. Based on the resource specifications we specified, the scheduler chooses the workernode of enough capacity to run the pod on the cluster.

There are 2 types of metrics we can specify

1. cpu = how many number of cpus to be allocated
2. memory = how much amount of memory to be assigned

These limits will be specified to the individual container of a pod, not at the pod level. so that the sum of individual container resource specifications would be considered as the final resource limits in scheduling and running a pod on the cluster

There are 2 types of limits are there

1. requests = requests indicates the initial capacity to be allocated in running the pod
2. limits = the max capacity a pod can be give for execution

In case if we have not specified the "limits" metrics then it considers the max cpu/memory the memory to be allocated as

1. no limit, and allocate how much ever has been requested
2. while creating a namespace, the administrator can set a max default limit of the resources to be allocated to the pods, which would be applied for all the

containers running inside the namespace, given if limits not specified at the pod level

In case if we have not specified the "requests" resource spec for cpu/memory of a pod, then it would be considered by default same as "limits" spec only

It is always advised to specify the resource declarations within the pod spec file to better manage the pod on the cluster

What are the different states in which a pod can exist in a Kubernetes cluster?

A pod in a Kubernetes cluster can exist in 5 different states, which is also referred to as "pod lifecycle"

1. pending = when we send a request for creating a pod, the api manager upon accepting the request would send the request to the scheduler for creating the pod, at this moment the pod state is reported as "pending"
2. running = at least one of the containers inside the pod has been started and readinessProbe on the pod has been passed, then the pod is reported as running
3. succeeded = all the containers within the pod have been exited with an exitcode as zero, then the pod is reported as succeeded
4. failed = when at least one of the containers inside the pod has been exited with non-zero exitcode then the pod is reported as failed
5. crashloopbackoff = when a pod is repeatedly failing for execution after a successive restarts, then to avoid further scheduling of the same pod for execution, Kubernetes marks the state of the pod as "crashloopbackoff" indicating the pod should not be scheduled for further execution, since it is repeatedly failing

Working with Labels and Annotations in Kubernetes

Labels:

Labels are arbitrary key/value pairs we can attach to a Kubernetes object, these are used for identifying and accessing the objects over the cluster. For a Kubernetes object we can assign any number of key/value pair labels but, they should appear only once and should be unique

We can define labels for the Kubernetes objects in 2 ways

1. we can declare the labels in the spec or manifest files directly
2. we can attach/detach the labels dynamically at runtime

As specified earlier we can attach labels for different types of objects like

1. pod
 2. replicaSet
 3. deploymentSet
 4. Service
- etc

#1 we can declare the labels in the spec or manifest files directly

How to declare labels in the pod specfile while creating the pod?


```

roadster-pod.yml
apiVersion: v1
kind: Pod
metadata:
  name: roadsterpod
  labels:
    app: roadster
    version: 1.0
    env: stage
spec:
  containers:
    - name: roadster
      image: techsriman/roadster:1.0
      ports:
        - name: tomcatport
          containerPort: 8080
          protocol: TCP

```

1. How to see all the labels attached to an object?

```
kubectl get pods --show-labels
```

2. we can query or search for objects based on the labels using -l switch?

```
kubectl get pods -l key=value
```

#2. how to dynamically attach labels to the object on the cluster?

```
kubectl label pods podName key=value
```

Annotations

Annotations are used for attaching arbitrary information which is non-identifier data to a Kubernetes object. These are only used as documentation helpers which we can read or use through the Kubernetes metadata API.

```

roadster-pod.yml
apiVersion: v1
kind: Pod
metadata:
  name: roadsterpod
  labels:
    app: roadster
    version: 1.0
    env: stage
  annotations:
    license: GPL License
    Warranty: product comes under limited warranty
spec:
  containers:
    - name: roadster
      image: techsriman/roadster:1.0
      ports:

```

```
- name: tomcatport
  containerPort: 8080
  protocol: TCP
```

we can see the annotations attached to an object using
`kubectl describe pod podName`

22-02-2023

configmaps- config-secret

ConfigMaps

What are ConfigMaps and what is the purpose of ConfigMaps?

Every software application uses configuration information pertaining to an external resource/system it is using to perform operations like an

1. database management system
 2. enterprise cache (redis, memcache)
 3. remote filesystem
 4. an external api provided by a vendor
- etc

How can the software application maintain or use this configuration within their application?

There are many ways of maintaining the configuration information

#1

The application can directly hardcode the configuration values within their programs in order to communicate and perform operations on the external resources/systems. Writing the configuration values directly within the sourcecode of the application is not recommended, because we run into a lot of problems as described below.

dis-advantages:

1. whenever there is a change in the configuration values, the developer has to modify the sourcecode, build/package and redeploy the application which takes a huge amount of time and rework for a configuration change
2. from one env to another env, the configuration values of the external resources would be different, so while moving the application across the env, we need to modify the sourcecode which again runs into the same problem we discussed above

#2

To avoid the above problem in maintaining the configuration information, we need to externalize the configuration values into an external file like properties, yaml or xml files. The developer has to write these configuration values within these configuration files and read those values in the programs to perform operations on the external resources.

advantages:

1. since the configuration values are placed in non-program files, a change in configuration needs to modify the configuration files which does not require rebuild/repackaging or redeploying the applications. So the configuration changes

can be easily reflected

2. for different environments we have different configuration values, so we can create multiple configuration files pertaining to each env and we can run the application against those configurations

dis-advantages:-

1. incase if multiple replicas of the application has been deployed across the nodes of the cluster for high availability and scalability, then maintaining these configurations and changing the configuration values across all the instances will be difficult

#3 centralize the configuration and distribute it across all the instances of the application

In addition it is advised to design the application to accept the configuration values from the environment variables, so that we can easily override the configurations while launching the application

How to run such applications that accepts the configuration values from an external source like env variables aspart of the kubernetes cluster?

23-02-2023

ConfigMaps

An Application has to be designed to read the configuration values as an input through environment variables so that while running the application the devops engineer can pass these values as an input by configuraing them as environment variables.

How to lauch or run an application which accepts the configuration interms of env variables on a kuberneters infrastructure?

The env variables should be seeded into the container, while launching the container so that these env variable values will be available as an input to the application. In case of kubernetes we are not creating/lauching the containers rather the controlplane takes care of creating the container, so we need to specify which what env/values with which the container application has to be launched by writing them in spec file

If we write these env variable with values in the pod specfile, there are few problems are there

1. everytime there is a change in the value, we need to modify the pod specfile, which is an unnecessary maintainance

2. the same env variables with values may have to be reused across different applications running on the cluster, so if we write the env variables with values locally in the pod spec file these will get duplicated across the applications so a change in these values inccurs huge efforts and time

Instead of writing these configuration values in pod spec file place them inside

the configmap object in kubernetes cluster. These configuration values we placed inside the ConfigMap can be accessed as inputs into the pod application in 3 ways

1. environment variables = we can pass these configMap values as environment variables into the pod application by refferring them in pod specfile
2. command-line arguments = we can pass these configMap values as commandLine-arguments while launching the application
3. Through ConfigMap api = The containerized application, the developer can write the code in reading the values from ConfigMap object stored on the kubernetes cluster (not recommended)

How to work with ConfigMap?

In order to use the configuration values from the ConfigMap, we need to create the ConfigMap object with keys/values on the cluster. For creating the ConfigMap object we need to write ConfigMap spec file as below.

```
coronaguidelines-configmap.yml
apiVersion: v1
kind: ConfigMap
metadata:
  name: guidelinesconfigmap
  labels:
    app: corona
data:
  oxygenLevels: 85
  quarantine: 20
  liters: 5
  temperatureLevels: 99 - 100
```

with these properties we created an ConfigMap object in kubernetes cluster, now we need to pass these properties in running the pod on the cluster as below

```
corona-pod.yml
apiVersion: v1
kind: Pod
metadata:
  name: coronapod
  labels:
    app: corona
spec:
  containers:
    - name: corona
      image: techsriman/corona:1.0
      ports:
        - name: tomcatport
          containerPort: 8080
          protocol: TCP
      env:
        - name: guidelines.oxygenLevels
          valueFrom:
            configMapKeyRef:
```

```

                                name: guidelinesconfigmap
                                key: oxygenLevels
- name: guidelines.liters
  valueFrom:
    configMapKeyRef:
      name: guidelinesconfigmap
      key: liters
- name: guidelines.temperatureLevels
  valueFrom:
    configMapKeyRef:
      name: guidelinesconfigmap
      key: temperatureLevels

```

```

=====
=====
27-02-23

```

What are ConfigMaps, why do we need to use them?

ConfigMap is a Kubernetes object in which we can store key/value pair data inside it and we can use the ConfigMap data inside the applications that are running in the pod containers.

There are 3 ways we can use the data inside the pod applications

1. environment variables / command-line arguments to the program
2. application programs can use Kubernetes APIs to access ConfigMap object that is stored on the cluster
3. We can mount ConfigMap with ConfigurationFile as volumes into the container through podspec

3. How to pass ConfigMap with configuration file as a volume in the podspec?

```

corona-configmap.yml
apiVersion: v1
kind: ConfigMap
metadata:
  name: coronaconfigmap
data:
  corona.properties |
    oxygenLeves=85
    liters=5
    quarantine=20
    temperatureLevels=99 - 102

```

Now while writing the podspec file we need to mount the properties file as a volume mount into the container, so that developers while building the application will have the logic for reading the file from the mountLocation

```

corona-pod.yml
apiVersion: v1
kind: Pod

```

```

metadata:
  name: coronapod
spec:
  container:
    - name: corona
      image: techsriman/corona:1.0
      ports:
        - name: tomcatport
          containerPort: 8081
          protocol: TCP
      volumeMounts:
        - name: coronavolume
          mountPath: /config
          readOnly: true

  volumes:
    - name: coronavolume
      configMap:
        name: coronaconfigmap
        items:
          - key: corona.properties
            path: "corona.properties"

```

here we defined an volume "coronavolume" populated with ConfigMap data (corona.properties) inside it. Then we are mounting the volume inside the container of a pod using volumeMounts under a specific directoryLocation /config

lets connect tommorow looks like it is taking time

```

=====
=====

```

02-03-2023

Config Secrets

Kubernetes secrets let us store and manage sensitive informaton like passwords, ssh keys, encryption keys etc that are required as an input by an software application. We can store these secrets/credentials directly aspart of podspec or in a configMap as well, but storing these secrets in pod spec or configMap makes them insecure. everyone can read the information/secrets we stored in podspec or configMap and can grab the access to the end systems. Instead it is recomend to store such sensitive information in ConfigSecret

Note:-

By default when we store the credential information in ConfigSecret it will not encrypt the data while storing, rather it encodes the data into Base64 encoding and will be stored. That means we can read the values back into plain-text format and

hence these are not by default secured.

So kubernetes ConfigSecrets are stored in HashiCorp Vault by integrating kubernetes with vendor Vaults

While storing the sensitive data within the ConfigSecret we can attach type information to help us identify what type of secret we are storing in ConfigSecret. It is not mandatory to attach type information while storing a ConfigSecret but it is recommended so that we can easily understand what type it is while accessing. By default while storing if we don't specify the type, it treats the type as "opaque"

Kubernetes has provided built-in secret types, we can use these secret types while defining our own secrets

1. opaque = arbitrary data
2. kubernetes.io/service-account-token = The service account token is an system secret or kubernetes secret
3. kubernetes.io/dockercfg = serialized format of docker config file
4. kubernetes.io/dockerconfigjson = serialized format of docker config json file
5. kubernetes.io/basic-auth = username/password
6. kubernetes.io/ssh-auth = ssh keys
7. kubernetes.io/tls = ssl keys or public/private encryption keys

we can use the ConfigSecret in 3 ways like ConfigMaps

1. we can pass the secrets as environment variables or commandline arguments
2. we can mount the ConfigSecret as files in the pod container
3. The kubelet process itself uses these secrets for connecting to the docker container registry for pulling the images

How to create an ConfigSecret object for storing the database username and password?

```
airtel2-configsecret.yml
apiVersion: v1
kind: Secret
metadata:
  name: airtel2dbconfigsecret
type: kubernetes.io/basic-auth
stringData:
  username: root
  password: root
```

```
airtel2-pod.yml
apiVersion: v1
kind: Pod
metadata:
  name: airtel2pod
spec:
```

```

containers:
  - name: airtel2
    image: techsriman/airtel2:1.0
    ports:
      - name: tomcatport
        containerPort: 8081
        protocol: TCP
    env:
      - name: "spring.datasource.username"
        valueFrom:
          secretKeyRef:
            name: airtel2dbconfigsecret
            key: username
      - name: "spring.datasource.password"
        valueFrom:
          secretKeyRef:
            name: airtel2dbconfigsecret
            key: password

```

=====

04-03-2023

There are 4 components are there in kubernetes cluster

1. Control Plane
2. Worker Nodes
3. Kubectl
4. etcd

Within the Control Plane there are 3 components are there

1. Api Manager
2. Scheduler
3. Controller Manager

There are 5 types of controllers are there

1. ReplicaSet Controller
2. DeploymentSet
3. DaemonSet
4. Job
5. Service

1. ReplicaSet Controller

Pod is the smallest entity within the kubernetes cluster where multiple containers are kept together and are executed within a pod. these container may want to share common resources like FileSystem or has a common lifecycle to be packaged and run inside one pod.

We can create a pod in kubernetes cluster through manifest or pod specfile. There are few characteristics of a pod are there

1. one pod manifest creates one pod on the cluster in running state
2. pod will not survive by crash = If a pod has been crashed due to any reason it will not be replaced with another pod

In addition:

3. if we want to run #10 pods out of the same pod specfile (with same containers inside it), we need to create #10 pod manifest files with different pod names and create the pods manually, which is an difficult job
4. upon creating the #10 pods, we need to monitor them and ensure always those are running, incase if any one of the replica of the pod has been crashed, we need to take care of replacing them within another pod

So managing the multiple replicas of a pod and replacing them incase of crash by monitoring them is difficult job, So to overcome this problem kubernetes has provided ReplicaSet Controller

ReplicaSet Controller:

A ReplicaSet Controller can be imagined as an Reconciliation loop, where the ReplicaSet controller loops through all the workernodes of the cluster to identify whether the desired number of Replicas of a pod are running on the cluster or not. if the desired number of replicas are not met, then the ReplicaSet Controller talks to the Scheduler in bringing up the pods on the cluster to meet the desired state. If already the desired number of replicas are met, it goes into monitoring state to see if any pods has crashed over the course of time, it can replace them with another pod

To create an ReplicaSet Controller we need to write an ReplicaSet Specfile similar to an pod spec. within the ReplicaSet manifest/spec we need to embedded the podspec specifying how many replicas of this pod should be runing. without a podspec a replicaset spec will be meaningless

Let us consider we have an #4 workernode cluster, in which we created an replicaSet controller with desired number of replicas as #3, the kubernetes control plane or ReplicaSet controller makes best effort in distributing the #3 replicas of a pod on #3 different workernodes of the cluster for high availability, but there is no guarantee that always the replicas are running on uniquenodes of the cluster

From the above we can understand we always write the podspec inside the replicaset spec with replicas to bring up desired number of pods on the cluster

```
sailor-replicaset.yml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: sailorreplicaset
  labels:
    app: sailor
spec:
  replicas: 2
```

```

        selector: [through which we specify which pods of the specified replicas to
be running]
        matchLabels:
            app: sailor
            version: 1.0
template:
    metadata:
        labels:
            app: sailor
            version: 1.0
    spec:
        containers:
            - name: sailor
              image: techsriman/sailor:1.0
              ports:
                - name: tomcatport
                  containerPort: 8080
                  protocol: TCP

```

```

=====
=====
05-03-2023

```

Deployment

Deployment is another way of deploying the application onto the kubernetes cluster and releasing it to the customers. Using Deployment controller any changes to the pod template can be rolled out in a controlled way

There are several deployment strategies supported by Deployment controller

1. recreate = terminate the old version and release the new one
2. ramped = release a new version on a rolling update fashion one after the another
3. blue/green = release a new version alongside of the old version then switch the traffic to newer version
4. canary = release a new version to a subset of users then proceed to a full rollout
5. a/b testing = release a new version to a subset of users in a precise way (HTTP headers, weight etc). A/B testing is a technique for making business decisions based on statistics. It doesn't work out of box with kubernetes, we need to do extra setup to implement it like (istio, traefik, nginx/haproxy etc)

Deployment controller applies the strategy of rolling or releasing the pod replicas on the cluster, by which we can understand without a ReplicaSet there is no deployment controller exists to manage the releases. So within the Deployment spec always we embed the ReplicaSet spec

speed-deployment.yml

```

-----
apiVersion: apps/v1
kind: Deployment
metadata:
  name: speeddeployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: speed
  strategy:
    type: Recreate
  selector:
    matchLabels:
      app: speed
  template:
    metadata:
      labels:
        app: speed
    spec:
      containers:
        - name: speedcontainer
          image: techsriman/speed:1.0
          ports:
            - name: tomcatport
              containerPort: 8080
              protocol: TCP

```

once we created the deployment with the above deployment spec, we can apply the spec to create deployment on the kubernetes

```
kubectl create -f speed-deployment.yml
```

#1. how can we see the deployments on the cluster?

```
kubectl get deployments
```

#2. describe deployments

```
kubectl describe deployment deploymentName
```

when we create an deployment, kubernetes internally creates an replicaSet in rolling out the pods based on the deployment

```

-----
How to rollout the newer version of the application?

```

1. make the changes in the application, and have these changed build, tested and packaged and distributed it as docker image into the docker container repository the image name should be bumped up when we are releasing the newer version of the application

2. upon the new version of the image is published and ready, we need edit the deployment we have already created for that application on the cluster

There are 2 ways we can modify the deployment

1. using kubectl set command
2. updating the specfile on the cluster

1.

```
kubectl set image deployment/deploymentName container=newImage:newVersion
```

for eg..

```
kubectl set image deployment/speeddeployment speedcontainer=techsriman/speeddep:2.0
```

2. run the below command that opens the existing deployment spec on the cluster

```
kubectl edit deployment deploymentName
```

```
kubectl edit deployment speeddeployment
```

#3. how to see the revision history of the rollout changes?

revision history means the modifications done on a deployment, we can see them using the below command

```
kubectl rollout history deployment/deploymentName
```

everytime when we modify the deployment, kubernetes internally creates a new replicaSet controller with the latest changes and modifies the old replicaSet replicas to 0 to rollout the newer pods. the way the old and new replicaSet controller will be changes is depends on the strategy we choosen.

strategies:

1. Recreate

The recreate strategy means terminate all the running instances then recreate them with newer version

spec:

```
replicas: 2
strategy:
  type: Recreate
```

advantages:-

1. application state entirely renewed
2. no need of additional infrastructure to be created or planned to rollout the newer version of the application
3. cost of releasing the newer version is less

dis-advantages:-

1. downtime and the amount of time the application will be un-available depends on the time it takes to boot the instances

=====

06-03-2023

Deployment Controller

Deployment controller is another way of deploying the application onto the kubernetes cluster. Through deployment controller any changes in the pod template can be rolled out in a controlled way

For eg.. if we have an pod application of version: v1 running on the kubernetes cluster, and we have an updated version of the application: v2 is available. To update the existing application of v1 version on cluster with v2 version the Deployment controller helps us in releasing by providing various different deployment strategies. There are 5 deployment strategies are supported by Deployment Controller

1. Recreate
2. Ramped
3. blue/green
4. canary
5. a/b testing

#1. Recreate

In Recreate the existing pods on the cluster will be terminated and the new version of the pods will be rolled onto the cluster.

advantage:-

1. the new version can be rolled out at one shot
2. no need of additional infrastructure to be planned for making an release

advantage:-

1. downtime, and the downtime depends on how long the new version of the application takes time to boot up

#2. Ramped [Slow Rollout] [Rolling Update]

A Ramped deployment updates the pods in a rolling update fashion, a second replicaSet will be created with the new version of the application, then the number of replicas of the old version is decreased and the new version is increased until the desired number replicas are reached.

upon creating 2 replicaSet controllers, it will reduce the number of replicas of the old ReplicaSet to 1 less than the total for eg.. if it is 3 then it makes as 2. Then updates the new ReplicaSet controller with replicas = 1 upon the new version of the pod readinessCheck has been passed again it goes to old ReplicaSet and decreases the replicas by 1 and increments the replicas for new ReplicaSet

```
speed-deployment.yml
appVersion: apps/v1
kind: Deployment
metadata:
```

```

    name: speeddeployment
spec:
  replicas: 3
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 2          # how many pods can be added at a time
      maxUnavailable: 0    # maxUnavailable pods can
existing during this rolling update

```

advantages:-

1. version is slowly released across the instances
2. no downtime of the application

dis-advantage:-

1. rollout/rollback can take more time
2. no control over the traffic
3. supporting multiple api versions is very hard

How to rollback to the specific version in the rollout history?

```
kubectl rollout undo deployment/deploymentName --to-revision=revisionNo
```

How to scale the deployment

```
kubectl scale deployment/deploymentName --replicas=4
```

#3. Blue/Green Deployment

Unlike the Ramped version where the pods are updated from old to the new version in rollout fashion, in case of blue/green deployment along with the blue version of the pods the green version of the pods is also available alongside.

After testing the new version of the pods, we update the Service object

For supporting blue/green deployment strategy, we need to create one more deployment spec to rollout the new version of the application and once it is been tested, the re-point the service to point to the latest pods and delete/scale-down the older deployment

advantages:-

1. instant rollout/rollback
2. zero downtime for the application
3. avoids versioning issues, because it changes the entire cluster state at one go

dis-advantage:-

1. requires more infrastructure for every release
2. handling stateful applications will be hard

=====

=====

07-03-2023

how to pause and resume the deployment?

if we want to perform multiple updates/changes to the pod, and wanted to deploy all of these changes at once then we need to pause and resume the deployment

1. Pause

kubectl rollout pause deployment/deploymentName

upon applying the changes we can resume the deployment using the below command

kubectl rollout resume deployment/deploymentName

1. history

2. scale

3. to-revision

4. rollout undo

5. pause

6. resume

out of box kubernetes deployment controller supports only 2 deployment strategies

1. Recreate

2. RollingUpdate

For realtime application deployments, these 2 strategies are not sufficient, because each of them has their own downsides. So people use different strategies of their own in rolling out the changes to the pod application on to the cluster

3. blue/green = release a new version of the pod applications alongside the old versions and upon completing the testing, switch the traffic to new version and obsolete the older application

4. canary = the canary strategy is also similar to blue/green only, the only difference is we release the new version of the application to the subset of users, then based on feedback we rollout the full version

5. a/b testing = release a new version to the subset of users in precise way / controlled way. For eg if the user is sending the request with a specific request header or cookie, then send the user to new version of the application, other let the user access to access the older version

Service

Service is adding networking to the pods that are running on the kubernetes cluster. By default when we create a pod on the node of a cluster, the pod has few characteristics:

1. The pod will be accessible within the Node of the cluster and cannot be accessed by any other pods that are running on other nodes of the cluster

2. The pod will be assigned with ephemeral ip address, which would be renewed upon

a pod restart

but we wanted the pod applications to be

1. accessed across the nodes of the cluster, so that other pods can access our pod
2. we want the pod applications to be exposed over a fixed ip address so that in case of a pod crash, the other applications using our pod will not be impacted upon recovery
3. we want to loadbalance the requests that are received by multiple instances of a pod
4. we want the pod applications to be exposed and made them accessible to the public world

all these can be achieved through Service

There are 5 types of Services are supported by Kubernetes Cluster

1. ClusterIP
2. NodePort
3. Loadbalancer
4. Ingress
5. Headless Service

=====

08-03-2023

Service

Service is adding an network to the pods of the kubernetes cluster. By default a pod is accessible within the node on which it has been created and has few more characteristics

1. A pod will not be recovered upon crash
 2. ephemeral ip address is assigned to the pod, so that upon crash and recovery a new ip address will be assigned
- but we wanted the pod to be
1. accessed by other pods on the cluster without worrying about the ip address being changed
 2. pod should be exposed to the external world
 3. the traffic should be load balanced across the pod replicas
- we can achieve all these things through the help of Service

There are 5 types of services are there

1. ClusterIP
2. NodePort
3. Loadbalancer
4. Ingress
5. Headless Service

#1. ClusterIP Service

whenever we create a service in kubernetes cluster, by default the service type is cluster ip if we don't specify any type. If we want our pod application to be exposed within the cluster and make it accessible over fixed ip and loadbalance the

traffic we need to use ClusterIP Service.

The clusterIP Service will be assigned with an ip address within the cluster range, using which the other pods of the cluster can access the pod.

usually we expose backend applications like microservices to the frontend applications or database pods to the java applications on the cluster using ClusterIP Service.

#2. NodePort

NodePort service the name itself indicates we open a Port on each WorkerNode through which we receive traffic and forward the to the targetPort on which the pod application is running. We can directly access a Pod running on a workernode using the fixed ip and port number using NodePort Service.

It is not meant for loadbalancing and distributing the traffic to the pods that are running across the nodes of the cluster.

Usually it is not recommended to use NodePort Service, since it exposes the Port of a WorkerNode over fixed ip address that creates a security breach. The purpose of NodePort is to directly access the application that is running on a workernode during development/testing phases only

Note: To have a pod application running on a worker node to be directly accessible to the public world, the WorkerNode should have public ip address.

=====

09-03-2023

NodePort

NodePort is used for exposing a pod application that is running on a workernode of the cluster directly to the external world over a fixed port. The name itself tells us, it opens a port on the workernode through which it makes the underlying pod application running on the node to be made accessible to the world.

=====

10-03-2023

NodePort

NodePort Service is used for exposing a pod application that is running on a workernode directly to the external world. NodePort opens a port on the workernode (within the range: 30000 - 32767) and forwards the request to the targetPort on which the Pod application is running.

When we create a NodePort Service, The NodePort will be created across all the Nodes of the cluster opening a port by itself. The traffic received on the WorkerNode NodePort will be forwarded to NodePort Service Port, through which be

forwarded to the targetPort on which the application is running on a pod

There are 3 Ports are opened in NodePort Service

1. NodePort = opens a port on the WorkerNode allowing the external network traffic
2. Port = Service Port to which the Traffic will be forwarded upon receiving the external traffic on the NodePort, now the service does the discovery of the pods based on selector to forward the request to the targetPort
3. TargetPort = port on which the pod applications are running

The default username/password for login into minikube server is

docker/tcuser

ssh docker@minikube

Persistence Volume and Persistent Volume Claim

How to run stateful applications on kubernetes pods?

The Stateful applications generates the data during their execution, that is by default written onto the container write layer of the container of the pod in which it is running. In case of a crash, the data written by the container will not be survived.

How can we persist the data that is generated by the pod application that is running inside the container of a Pod, so that in the event of crash the data would be retained?

That is where kubernetes has introduced Persistent Volume and Persistent Volume Claim

Persistent Volume = is a storage defined on the kubernetes cluster and it is created and maintained by the kubernetes administrator

Persistent Volume Claim = is a request for that storage to be consumed by the devops engineer that can be used as part of a pod application

There are few attributes we need to define while creating a persistent volume:

1. storageClassName = indicates the type of storage to be created on the cluster
2. accessMode:

ReadOnly = only the pods can read the data from this volume

ReadWriteOnce = only one pod can write the read/write the data at one time

ReadWriteMultiple = Multiple pods are allowed to read/write at the same time

3. capacity = storage size to be assigned for that persistent volume

How to define a persistent volume?

pv.yml

apiVersion: v1

kind: PersistentVolume

metadata:

name: pv1

spec:

storageClassName: pvClass

capacity:

```
        storage: 2Gi
    accessMode:
      - ReadWriteMultiple
    hostPath:
      path: /u01/data
```

pvc.yml

```
-----
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc1
spec:
  storageClassName: pvClass
  accessMode:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
```

speed-pod.yml

```
-----
apiVersion: v1
Kind: Pod
metadata:
  name: speedpod
spec:
  containers:
    - name: speedcontainer
      image: techsriman/speeddep:2.0
      ports:
        - name: tomcatport
          containerPort: 8080
          protocol: TCP
      volumeMounts:
        - name: speedvolume
          mountPath: /u01/app

  volumes:
    - name: speedvolume
      persistentVolumeClaim:
        claimName: pvc1
```

```
=====
=====
11-03-2023
```

Persistent Volume and Persistent Volume claims

Persistent Volume = it is a storage that is defined on the kubernetes cluster and

it is created by the kubernetes administrator reserving the storage location on the worker nodes

Persistent Volume Claim = is a request for that storage to be consumed by the devops engineer aspart of the pod application

There are 3 attributes we need define to create an persistent volume

1. storageClassName
2. accessMode
3. capacity

mysql-pv.yml

```
-----
apiVersion: v1
kind: PersistentVolume
metadata:
  name: mysqlpv
spec:
  storageClassName: mysqlStorageClass
  capacity:
    storage: 2Gi
  accessMode:
    -ReadWriteMultiple
  hostPath:
    path: /u01/data #location on the workernode the volume should be
created
```

to use the part of the storage we need to create an pvc

mysql-pvc.yml

```
-----
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mysqlpvc
spec:
  storageClassName: mysqlStorageClass
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
```

mysql-pod.yml

```
-----
apiVersion: v1
kind: Pod
metadata:
  name: mysqlpod
spec:
```

```

    container:
      - name: mysqlcontainer
        image: mysql:8.0
        ports:
          - name: mysqlport
            containerPort: 3306
            protocol: TCP
        volumeMounts:
          - name: mysqlvolume
            mountPath: /u01/mysql
    volumes:
      - name: mysqlvolume
        persistentVolumeClaim:
          claimName: mysqlpvc

```

```

=====
=====

```

13-03-2023

Ingress

Ingress is an another type of Service in kubernetes, that is used for exposing the pod applications to the external world.

Ingress is an another controller of the kubernetes, that receives the request from the external world and routes it to the Service component.

by default withan kubernetes or minikube install Ingress controller will not be available, we need to enable the ingress explicitly
minikube addons enable ingress

There are different ingress controller component providers are there in the market one such provider is Nginix

These are internally httpd servers which receives the request over an domainName and proxy the request to the backend

now we need to write an ingressService that receives request over a domain or host: covido.org

then forward the request to the ClusterIP Service (clusterip)

covido-service.yml

apiVersion: v1

kind: Service

metadata:

name: covidoclusteripservice

spec:

type: ClusterIP

selector:

app: covido

version: v1

ports:

- port: 8080
target: 8080

covido-ingress.yml

apiVersion: networking.k8s.io/v1

kind: Ingress

metadata:

name: covidoingress

annotations:

nginx.ingress.kubernetes.io/rewrite-target /\$1

spec:

rules:

- host: covido.org

http:

paths:

- path: /

pathType: Prefix

backend:

service:

name:

covidoclusteripservice

port:

number:

8080

http://covido.org/index

=====

15-03-2023

Job

Job is used for performing an operation on the kubernetes cluster. A Job creates one or more Pods and executes until the specified number of executions are successful

numbers-job.yml

apiVersion: batch/v1

kind: Job

metadata:

name: numbersJob

spec:

template:

metadata:

name: numberspod

spec:

containers:

```

        - name: numberscontainer
          image: ubuntu:20.03
          command:
            - "bin/bash"
            - "-c"
            - "for i in 1 2 3 4 5 6 8 9 0 ; do
echo $i ; done"
          restartPolicy: Never

```

DaemonSet

nginx-daemonset.yml

```

apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: nginxdaemonset
spec:
  selector:
    matchLabels:
      name: nginx
  template:
    metadata:
      labels:
        name: nginx
    spec:
      containers:
        - name: nginx
          image: nginx

```

=====

19-03-2023

Premetheus can be installed on minikube using helm charts

helm repo add prometheus-community

<https://prometheus-community.github.io/helm-charts>

helm repo update

helm install prometheus prometheus-community/prometheus

then expose the prometheus server over nodeport

kubectl expose service prometheus-server --type=NodePort --target-port=9090

--name=prometheus-server-ext

service/prometheus-server-ext exposed

to access the prometheus dashboard run the below command

```
minikube service prometheus-server-ext
```


